

C++ Memory management

- Storage types
- Allocation and deallocation
- *malloc()/free()* vs. *new/delete*
- Memory management internals

C++ Memory Management – Introduction

- C++ added necessary memory management support for object-oriented programming
- C++ fixed some loopholes and enhanced the memory management compared to plain C compilers
- This chapter shows:
 - The memory management model in C++
 - The three types of data storage
 - The various versions of operators (*new* and *delete*)
 - Some technique and guidelines for effective and bug-free storage usage

C++ Memory Management – Types of storage

- C++ has three fundamental types of storage:
 - Automatic storage
 - stack memory
 - Static storage
 - static data
 - Free storage
 - dynamically allocated data
- Each of them has different semantics of object initialization and lifetime

C++ Memory Management – Types of storage – Automatic storage

- The automatic storage (also called *stack storage*) is used for
 - Local objects that are not explicitly declared as *static* or *extern*
 - Local objects that are declared *auto* (default) or *register*
 - Function arguments
- *Created automatically* on the stack upon entering a function or a block
- *Destroyed automatically* when the function or the block exits
- At entrance, a *new copy* is always created
- **The default** value of automatic variables and non-class objects is **indeterminate**

C++ Memory Management – Types of storage – Static storage (1)

- The static storage is used for
 - Global objects
 - Static data members of a class
 - Namespace variables
 - Static variables in functions
- The address of a static object **remains the same** throughout the program's execution cycle
 - Constructed only once during the lifetime of the program
- **By default**, initialized to **zero**
 - Initialized by its constructor, if needed

C++ Memory Management – Types of storage – Static storage (2)

```
int num; // global variables have static storage
```

```
int func( void ){  
    static int calls; // initialized to 0 by default  
    return ++calls;  
}
```

```
class C {  
private:  
    static bool b; // b has static storage  
};
```

```
namespace NS{  
    std::string str; // str has static storage  
}
```

C++ Memory Management – Types of storage – Free storage

- The free storage (also called *heap memory* or *dynamic memory*) contains
 - Objects created by the *new* operator
 - Variables created by the *new* operator
- They persist until they are destroyed with the *delete* operator
 - Unreleased memory is not automatically returned to the operating system!
 - This produce **memory leaks**
- The address of a *free store* object is determined **at runtime only**
- The initial value of raw storage that is allocated with the *new* operator is **unspecified**



C++ Memory Management – Allocation and Deallocation Functions (1)

- C++ offers the following global functions for allocating and deallocating
 - *new* and *new[]*
 - *delete* and *delete[]*
- These functions are accessible from the header `<new>`
 - The inclusion of this header is not necessary, done implicitly

- The declarations look like:

```
void* operator new( std::size_t ) throw( std::bad_alloc );  
void* operator new[]( std::size_t ) throw( std::bad_alloc );  
void operator delete( void* ) throw();  
void operator delete[]( void* ) throw();
```


C++ Memory Management – Allocation and Deallocation Functions (2)

- The implicitly inclusion of *new/delete* does not implicitly include *std*, *std::bad_alloc* and *std::size_t*
 - The usage of these names needs an explicit inclusion of `<new>` header file

```
#include <new>
using namespace std;

char * allocate (size_t bytes);

int main
{
    char * buff = allocate( sizeof(char) );
    return 0;
}
```

C++ Memory Management – Semantics of Allocation and Deallocation (1)

- The allocation function returns an pointer to ***void****
- The first argument of an allocation function is of type ***std::size_t***
 - This corresponds to the requested memory size
- The allocation tries to allocate the requested memory block from the *free store* memory
- If the allocation was successful then it returns a pointer to the start of the reserved block
 - Failure will be discussed later

```
void* operator new( std::size_t ) throw( std::bad_alloc );
```

C++ Memory Management – Semantics of Allocation and Deallocation (2)

- The deallocation function returns nothing (*void*)
- The first argument is of type ***void****.
 - Additional arguments are possible
 - If the first argument is a pointer to *NULL*, then nothing is deallocated
 - The first argument must be a pointer that has been returned by an *new* operator
- Allocation and deallocation must be used in pairs

```
void operator delete( void* ) throw();
```

C++ Memory Management – *malloc()* / *free()* vs. *new* / *delete*

- C++ still supports *malloc()* and *free()* to ensure backward-compatibility
 - Combining legacy C code with C++ code
 - Write C++ code that is meant to be supported in C environment
 - Making *new* and *delete* implementable by using *malloc()* and *free()*
- Try to avoid using *malloc()* and *free()* in C++ code
 - They do not support object semantics
- *New* and *delete* are also **significantly safer**

C++ Memory Management – *malloc()* / *free()* vs. *new* / *delete* – Object Semantic Support

- *new* and *delete* automatically constructs respectively destructs objects
- *malloc()* and *free()* **only** allocate respectively deallocate memory space from the *heap*
 - The constructor and destructor **won't be invoked**

```
class C {...}
```

```
C* myNew( void ) {                               // very bad
    C* pC = static_cast<C*>(malloc(sizeof(C)));
    return pC;
}
```

```
C* cPt;
cPt = myNew();
```

C++ Memory Management – *malloc()* / *free()* vs. *new* / *delete* – Safety

- The *new* operator calculates automatically the size of the object that it constructs
 - With *malloc()* the programmer has to give this size
 - *malloc()* returns a pointer to *void**
 - needs an explicit *type casting*
- The *new* operator bypasses these two problems

```
int* p = static_cast<int*> malloc(sizeof(int));  
int* p2 = new int;
```

C++ Memory Management – *malloc()* / *free()* vs. *new* / *delete* – Extensibility

- The *new* operator can be overloaded
 - Special classes can thus implement its own *new* operator
 - This is not possible with *malloc()*
- Do not intermix *new* with *free()*, respectively *malloc()* with *delete*
 - **The result is undefined**
 - It is even possible, that *new* and *malloc()* use different heaps

C++ Memory Management – *new* / *delete* – Arrays (1)

- **new []** allocates an array of objects of a specific type
 - The returned value is a pointer to the first element of the array

```
int *p = new int[10];
bool equal = (p == &p[0]); // true
delete[] p;
```
- Object arrays must be delete with **delete []**
 - Using a plain *delete* ends in a unspecified behavior
 - *new []* stores the number of elements in the allocated array in a special way
 - The *delete []* retrives this number of elements and frees the correct number of elements
 - The correct number of destructors can be invoked
- Mismatch in *new []* and *delete []* causes memory leaks, heap corruptions or program crashes

C++ Memory Management – *new* / *delete* – Arrays (2)

- These rules apply also to arrays of fundamental types!
 - *delete []* doesn't invoke a destructor, but it still has to retrieve the number of elements in the array

```
void f()
{
    char* pc = new char[100];
    string* ps = new std::string[100];

    //...
    delete[] pc; // no destructors invoked

    delete[] ps // each member's destructor
                // is called
}
```

C++ Memory Management – Exceptions and *operator new* (1)

- Early C++ standards returned *NULL* if a *new* operation failed (e.g. not enough memory)
 - Same behaviour as *malloc()*
 - Applied also for *new []*
 - Check needed against *NULL*
 - Tedious and error prone
 - Long and time gourmand especially for arrays

```
char* p = new char [size];  
if( p == NULL ) // this was fine until 1994
```
- Failures in dynamic memory allocation are quite rare
 - Indicate an unstable system
- The *NULL* returning in case of problem was replaced with a thrown exception ***std::bad_alloc***

C++ Memory Management – Exceptions and *operator new* (2)

- Programs calling *new* directly or indirectly (e.g. STL) must have an exception handler for *std::bad_alloc*
 - Otherwise, the program will terminate with an unhandled exception
 - Testing the returned pointer against *NULL* makes thus no sense at all, because the exception arrives before a possible test
 - The pointer check only uses system resources in case of normal behavior

```
void f(int size) {  
    char* p = new char [size];  
    // ...use p safely  
    delete [] p;  
    return;  
}
```

C++ Memory Management – Exceptions and *operator new* (3)

- The usage of new with exception handling

```
#include <stdexcept>
#include <iostream>

using namespace std;
const int BUF_SIZE = 1048576L;
int main(){
    try{
        f(BUF_SIZE);
    }
    catch( bad_alloc& ex ){
        cout << ex.what() << endl;
        // ...other diagnostics and remedies
        return -1;
    }
    return 0;
}
```

C++ Memory Management – Exception-Free *operator new* (1)

- Sometimes it is undesirable to throw exception during the *new* operator
 - Exceptions are turned off for performance enhancement
 - Platform doesn't support exceptions
- Thus, the C++ standard has also a *new* operator, which **does not** throw an exception in case of failure
 - It returns a pointer to *NULL* in case of failure
 - This version of *new* operator takes an additional argument of type *const std::nothrow_t*
 - It exists for normal and for array *new* operators

C++ Memory Management – Exception-Free *operator new* (2)

- The usage will be

```
void f( int size ){
    char* p = new(nothrow) char [size];
    if( p == 0 ) // test against NULL
    {
        // ...use p
        delete [] p;
    }

    string* pstr = new(nothrow) string;
    if( pstr == 0 ) // test against NULL
    {
        // ...use pstr
        delete pstr;
    }
    return;
}
```

C++ Memory Management – Exception-Free *operator new* (3)

- The argument *nothrow* is defined and created in header `<new>`
 - `extern const nothrow_t nothrow;`
- The structure of *nothrow_t* looks like this
 - `struct nothrow_t {};`
- The empty structure won't be used to transport any information about the failure

C++ Memory Management – Placement *operator new* (1)

- Another version of the *new* operator exist, which allows to construct an object at a predetermined memory position
 - Building custom-made memory pools
 - Garbage collection
 - Mission critical applications
 - The memory is already allocated
- Construction of object with the placement *new* operator is faster
 - The construction doesn't need to allocate memory (only constructors are called)
 - **The memory must have been allocated in advance**

}

}

C++ Memory Management – Placement *operator new* (2)

- Destructors of objects constructed with the *placement new operator* has to be destructed **explicitly**
 - The `delete [] p` will delete only the char array but won't implicitly invoke the C++ classe's destructor

```
class C{
public:
    C() { cout<< "constructed" <<endl; };
    ~C(){ cout<< "destroyed" <<endl; };
};

int main(){
    char* p = new char [ sizeof(C) ];
    C* pc = new(p) C; // placement new
    //... used pc
    pc->C::~~C();
    delete [] p;
}
```

C++ Memory Management – Exceptions during Object Construction (1)

- The *new operator* performs two operations:
 - Allocation of memory from the free storage
 - Calls an allocation function
 - Construct the object on the just allocated memory
- What happens if the object construction fails (exception during the construction)?
 - Does the allocated memory consist?
- The allocated memory will be directly freed before propagating the exception
 - No memory leak is produced

C++ Memory Management – Exceptions during Object Construction (2)

```
#include <new>
using namespace std;

class C{/*...*/};

void __new() throw (bad_alloc){
    C* p = reinterpret_cast<C*> (new char [sizeof(C)]); // step
    1: allocate raw memory
    try
    {
        new(p) C; // step 2: construct the objects on previously
        allocated buffer
    }
    catch(...) // catch any exception thrown from C's
               constructor
    {
        delete[] p; // free the allocated buffer
        throw; // re-throw the exception of C's constructor
    }
}
```

C++ Memory Management – Alignment Considerations

- Pointers returned by the *new operator* have the correct alignment to be converted into another pointer

```
char* pc = new char[ sizeof( Employee ) ];
Employee* pemp = new(pc) Employee;
//...use pemp
pemp->Employee::~~Employee();
delete [] pc;
```
- This doesn't work with buffers that are allocated on the stack (automatic storage)

```
char pbuff [ sizeof( Employee ) ];
Employee* p = new(pbuff ) Employee; // undefined
behavior
```
- Neither it works on the previously *static* objects

C++ Memory Management – Alignment Considerations – Member Alignment

- The size of a *struct* or *class* might be bigger than the result of adding all internal data members
 - The compiler can add *padding* to realign members whose size does not fit exactly into a machine word
- To get the real size, use *sizeof()*

C++ Memory Management – Size of a Complete Object

- The size of a complete object can **never be** zero

```
class Empty{};  
Empty e;           // e occupies at least 1 byte
```

- The object e occupies at least 1 byte
- The compiler does not allow objects with zero bytes
 - With zero bytes, addresses of different empty objects could overlap
 - The compiler guarantees that each empty object has also an unique address
- Incomplete objects can have a size of zero byte
 - Subobject in a derived class

C++ Memory Management – *new* and *delete* declared in a namespace

- The operators *new* and *delete* can be declared in a **class scope**
- It is **illegal** to declare them in a *namespace*

```
char* pc;
```

```
namespace A{  
    void* operator new ( size_t );  
    void operator delete ( void * );  
    void func (){  
        pc = new char ( 'a' );  
    }  
}
```

```
void f() { delete pc; } // A::delete or ::delete?
```

C++ Memory Management – Overloading *new* and *delete* in a class (1)

- Overloading the operators *new* and *delete* for a given class *C* is possible

- The following statement would invoke the class' defined operators

```
C* p = new C;  
//....  
delete p;
```

- Class-specified versions of *new* and *delete* is useful when the default memory management is unsuitable
 - Different behavior in case of an allocation failure (see next slide)
 - Custom memory pool

C++ Memory Management – Overloading *new* and *delete* in a class (2)

```
class C{
    C() { cout << "constructed" << endl; };
    ~C() { cout << "destroyed" << endl; };
    void* operator new( size_t size ); // implicitly static
    void operator delete( void *p );  // implicitly static
};

void* C::operator new( size_t size ) throw( const char* ){
    void* p = malloc( size );
    if( p == 0 ) throw "allocation failure"; // instead of
                                              // std::bad_alloc

    return p;
}

void C::operator delete( void* p ){ free(p); }

int main(){
    try{ C* p = new C; delete p;}
    catch( const char* err ){ cout << err << endl; }
    return 0;
}
```

C++ Memory Management – Guidelines for Effective Memory Usage (1)

- Choosing the right type of memory storage is an important programmer's task
 - Impact of performance
 - Impact of security and reliability
 - Maintenance
- Prefer, if possible, **Automatic storage** to Free storage
 - Runtime overhead is little
 - Free storage needs interaction with the operating system
 - Free storage can be fragmented, thus finding a free spot can take some time
 - Exception handling also adds runtime overhead
 - Maintenance
 - Dynamic allocation might fail. Additional code must check this
 - Safety
 - Memory leaks can occur if dynamic objects are deleted more than one time or not deleted at all

C++ Memory Management – Guidelines for Effective Memory Usage (2)

```
void f( void )
{
    string* p = new string;
    // ...use p
    if( p->empty() != false )
    {
        // ...do something
        return; // OOPS! memory leak: p was not deleted
    }
    else // string is empty
    {
        delete p;
        // ..do other stuff
    }
    delete p; // OOPS! p is deleted twice if
              // isEmpty==false
}
```

C++ Memory Management – Guidelines for Effective Memory Usage (3)

```
void f( void )
{
    string s;
    // ...use s
    if( s.empty() != false )
    {
        // ...do something
        return;
    }
    else
    {
        // ..do other stuff
    }
}
```

Advanced C++ for Java
Programmers

Memory Management

C++ Memory Management – Local Object Instantiation

- The correct syntax for a local object instantiation by using the **default constructor** is:

```
string str; // no parenthesis
```

- Attention, using parenthesis has a completely different meaning

```
string str(); // declaration of a function
```

C++ Memory Management – Zero as an Universal_INITIALIZER

- The literal 0 (*integer*) can be used as an universal initializer for any fundamental data type

```
void *p = 0; // zero is implicitly converted to void*  
float salary = 0; // 0 is cast to a float  
char name[10] = { 0 }; // 0 cast to a '\0'  
bool b = 0; // 0 cast to false  
void(*pf)(int) = 0; // pointer to a function  
int (C::*pm)() = 0; // pointer to a class member
```

C++ Memory Management – Always initialize pointers

- An uninitialized pointer has an **undetermined value**
 - This is a *wild pointer*
 - It is quasi impossible to test if a *wild pointer* is valid
 - Especially when passed as an argument
 - Only *NULL* test is possible
- Do not trust to the *NULL-initialization* provided by the compiler!

```
void func( char* p );  
int main(){  
    char* p; // dangerous: uninitialized  
    //...many lines of code; p left uninitialized  
    if( p ) // erroneously assuming that a non-null  
    {  
        func( p ); // func has no way of knowing  
    }  
    return 0;  
}
```

C++ Memory Management – Deleting a Pointer More than Once

- Applying an *delete* to an already deleted pointer has a undefined behavior
 - The pointer is undefined after the first delete
- Assigning a delete pointer to *NULL* avoids this danger
 - Applying the delete operator to a *NULL* pointer is harmless

```
if( ps->empty() )  
{  
    delete ps;  
    ps = NULL; // safety-guard:  
    // further deletions of ps will be harmless  
}
```


C++ Memory Management – Storage reallocation

- In the C language, the function *realloc()* exists for resizing existing buffers
- This functionality **does not** exist in the C++ language
- Two possibilities exist in C++ to do something similar:
 - Allocating a new buffer with the new size, then copy the old buffer into the new one and finally delete the old buffer
 - Not very elegant and error prone
 - Inefficient and tedious
 - The better way for object / arrays that change their size often is to use containers of the STL
 - They change the size dynamically

C++ Memory Management – Local Static Variables (1)

- Local Static Variables (not to be confused with static class members) are initialized to *zero*
 - They are created before program's outset
 - They are destructed after program's termination
- They are accessible only within the declaration scope
 - Useful for function state storage
- Using classes with static data member is a better choice for keeping *states*
 - Better flexibility
 - Class member replace the static variables
 - Member functions replace the global functions

C++ Memory Management – Local Static Variables (2)

- Every derived object that inherits such a member function also refers to the same instance of the local static variables
- Static variables are problematic in a multithreaded environment
 - They are shared
 - They need to use locks to access them

C++ Memory Management – Conclusion (1)

- C++ offers tremendous diversity for dynamic memory allocation
- Three types memory storage exist (automatic, static, free)
- The *new* and *delete* operators have a lot of different versions
 - Exceptions, Exception-free, placement
- Garbage collector does not exist in C++
 - Additional runtime overhead
 - Destructors are not invoked immediately
- Use automatic memory allocation
- Use STL for object that grows and shrinks dynamically

C++ Memory Management – Conclusion (2)

- The source of most of the bugs in a C++ program must be searched in the

memory management

- Avoid these bugs by using if possible *automatic* memory allocation
- Avoid these bugs with a proper memory allocation and deallocation